

The Imitation Game: Exploring Brand Impersonation Attacks on Social Media Platforms

Supplementary Material

Bhupendra Acharya
CISPA

Dario Lazzaro
University of Genoa

Efrén López-Morales
Texas A&M University-Corpus Christi

Adam Oest
PayPal Inc.

Muhammad Saad
PayPal Inc.

Antonio Emanuele Cinà
University of Genoa

Lea Schönherr
CISPA

Thorsten Holz
CISPA

1 Images and Posts Analysis

In this supplementary material, we provide further information on posts and images analysis, such as clustering hyperparameters, cluster sizes, and examples of scammers impersonating official brands.

Posts Clustering Hyperparameters

In the hyperparameterization process for UMAP, default values from the BERTopic library¹ were employed. Specifically, we configured UMAP with `n_neighbors=15`, `min_dist=0.0`, `n_components=5`, and cosine similarity. We then set the `random_state` variable to a fixed value of 42 to preserve the reproducibility of our code.

For HDBSCAN, we chose `min_cluster_size=10` and used the Euclidean metric for clustering. To refine the clustering outcome, we adjusted `min_samples=50` to reduce the resulting number of clusters. Additionally, the default BERTopic method for outlier reduction (`reduce_outliers`)² was applied to minimize the presence of outlier samples in the clustering results. Finally, we employed a standard evaluation metric, i.e., silhouette score and visual inspection of resulting clusters to assess the quality and validity of the clustering outcomes.

Next, we provide additional details on post-clustering related to the top 100 clusters and all clusters in our dataset. The top 100 clusters were manually inspected to provide a qualitative study of the top 10 attack categories. A tabular representation is shown in Table 1.

Top 100 Clusters. For the top 100 clusters, we identified no singletons, with a minimum and maximum cluster size of 72 and 3236, respectively.

Table 1: Post-clustering analysis of the 100 scammer clusters and all clusters. The tables report the minimum, maximum, median, and mean number of scammers in each cluster.

Clusters Detail	100 Clusters #	All Clusters #
Number of clusters	100	543
Number of Singletons	0	0
Min. Cluster Size	72	57
Max. Cluster Size	3236	9713
Median Cluster Size	363.0	370.0
Mean Cluster Size	557.3	593.6

All Clusters. We identified a total of 543 clusters with no singletons. The minimum and maximum cluster size was found to be 57 and 9713, respectively. The median cluster size was found to be 370.

Profile Images Inspection. During our analysis, we examined a total of 219,465 images sourced from X, Instagram, Telegram, and YouTube. These images were then compared against a set of 2,484 official logos. For each account, we established a threshold for similarity with the nearest logo. If the similarity fell below this threshold, we categorized the account as miscellaneous, indicating that it was not impersonating a specific brand. To maintain rigor in our analysis, we opted for a stringent threshold of 0.8 similarity between the account image and the official logo, thereby minimizing false associations. In Table 2, we demonstrate that by lowering this threshold, the number of accounts impersonating official logos increases. However, this adjustment also introduces false positives into our analysis. The lower similarity can be attributed to several factors, such as poor quality of image resolution, use of outdated logos, and alterations to the original logo.

¹<https://maartengr.github.io/BERTopic/index.html>

²https://maartengr.github.io/BERTopic/getting_started/outlier_reduction/outlier_reduction.html

Table 2: Logo similarity matching results at different acceptance thresholds.

	Instragram		Telegram		X		Youtube	
Acceptance Threshold	0.80	0.75	0.80	0.75	0.80	0.75	0.80	0.75
Accounts analyzed	36803		55758		102360		24544	
Accounts impersonating brands	1068	2235	4105	6735	4265	15825	38	79
Accounts miscellaneous	35735	34568	51653	49023	98095	86535	24506	24465




Figure 1: Examples of scammers impersonating official brands. The leftmost image represents the official brand and the username where as the remaining other images are impersonated images of the official brand.